


Sets are unordered collections of things

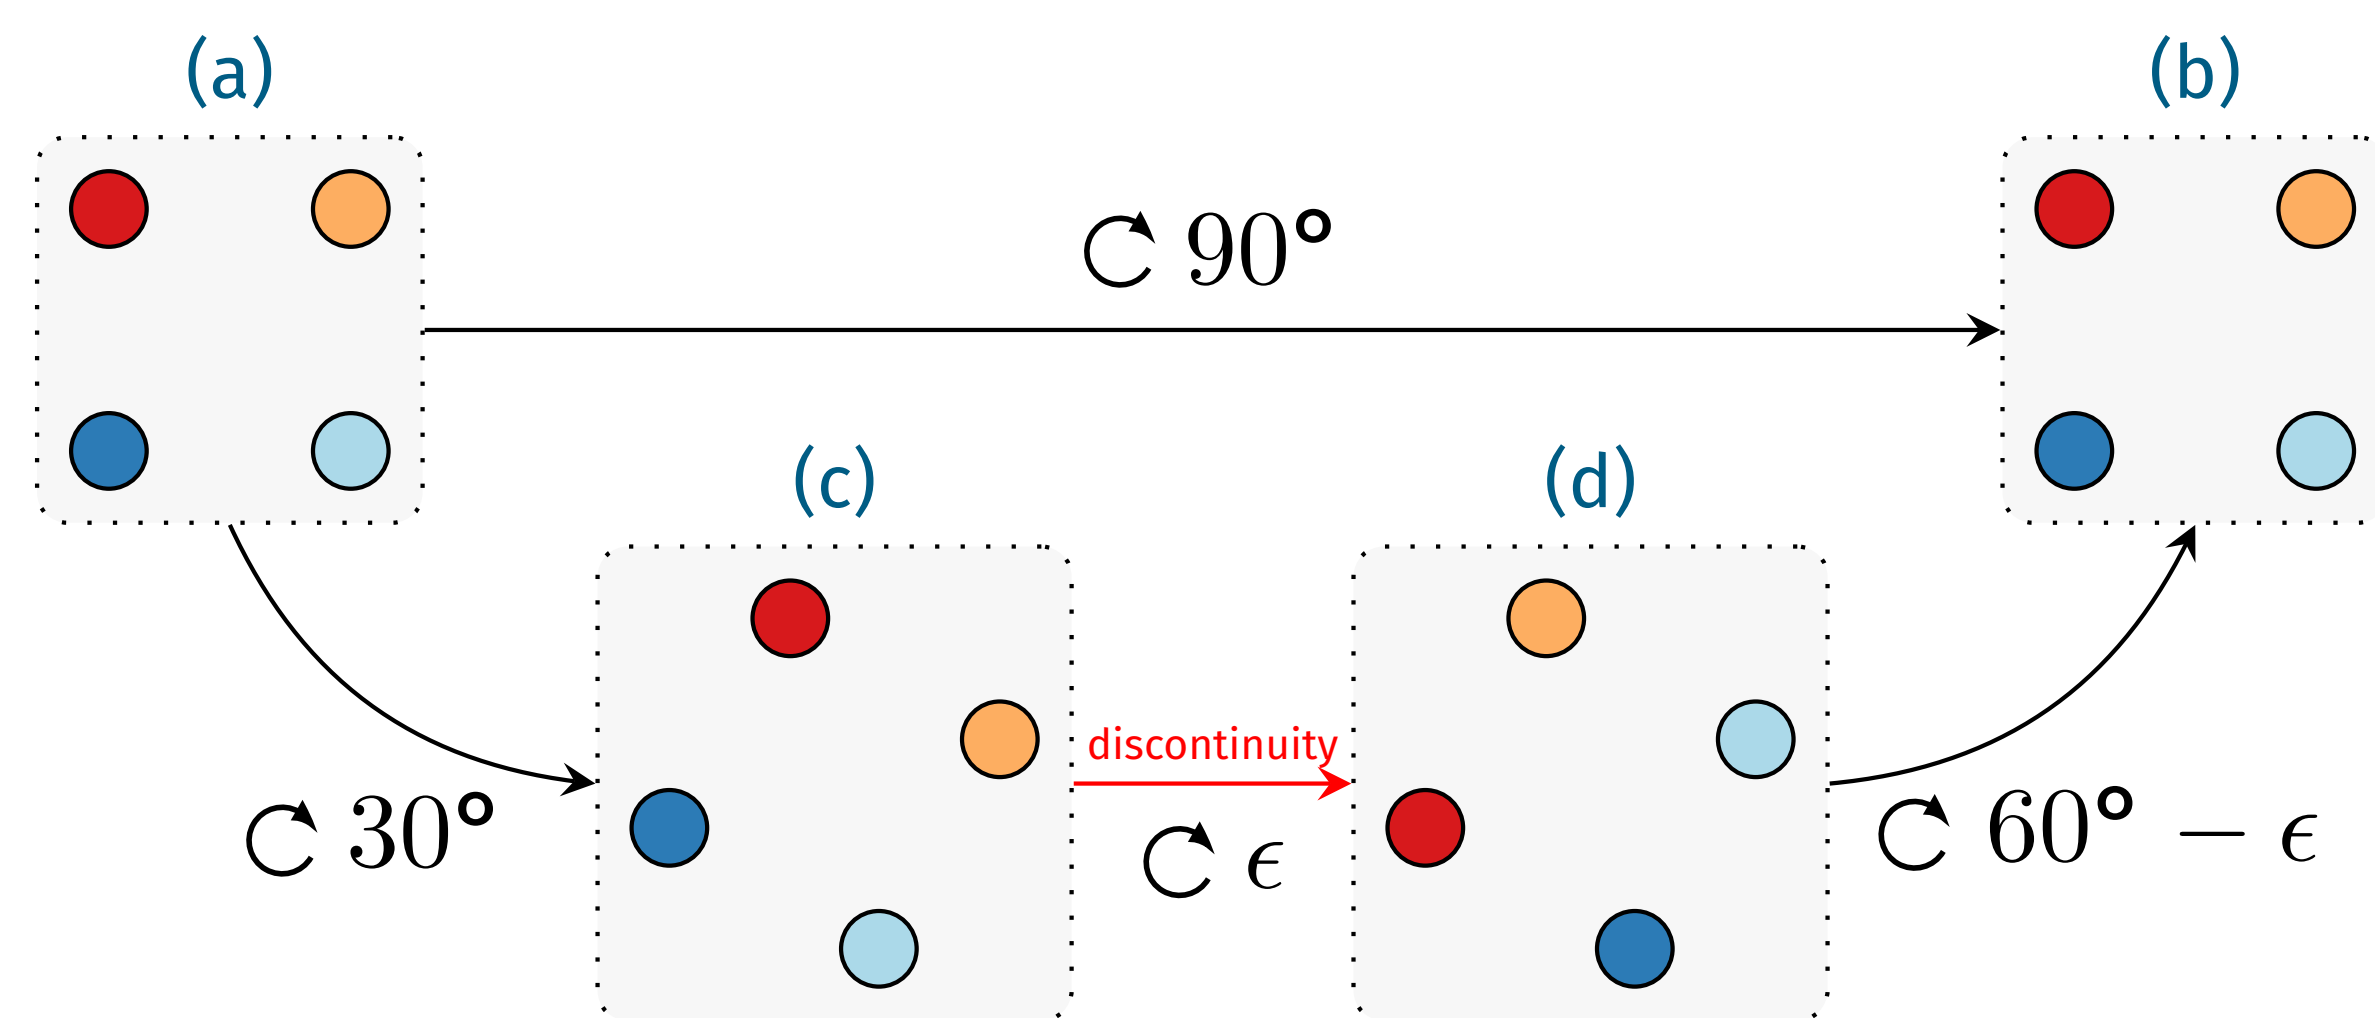
- Many things can be described as **sets of feature vectors**:
 - the set of objects in an image,
 - the set of points in a point cloud,
 - the set of nodes and edges in a graph,
 - the set of people reading this poster.
- Predicting sets means object detection, molecule generation, etc.
- This paper is about doing this **vector-to-set** mapping properly.
- Compared to normal object detection methods:
 - Anchor-free, fully end-to-end, no post-processing.

MLPs are not suited for sets

- Sets are **unordered**, but MLP and RNN outputs are **ordered**.
- **Discontinuities** from *responsibility problem*.
- Let's look at a normal set auto-encoder:



- The responsibility problem:



- (a) and (b) are the same set.
 - (a) and (b) encode to the same vector.
 - (a) and (b) have the same MLP output.
- (a) is turned into (b) by rotating 90°.
 - Rotation starts and ends with the same set.
 - MLP outputs can't just follow the 90° rotation!
 - There must be a **discontinuity** between (c) and (d)!
 - All the outputs have to jump 90° anti-clockwise.

Conclusion:

- Smooth change of set requires discontinuous change of MLP outputs.
- To predict **unordered sets**, we should use an **unordered model**.

To predict a set from a vector,
use gradient descent to find a
set that **encodes** to that vector.

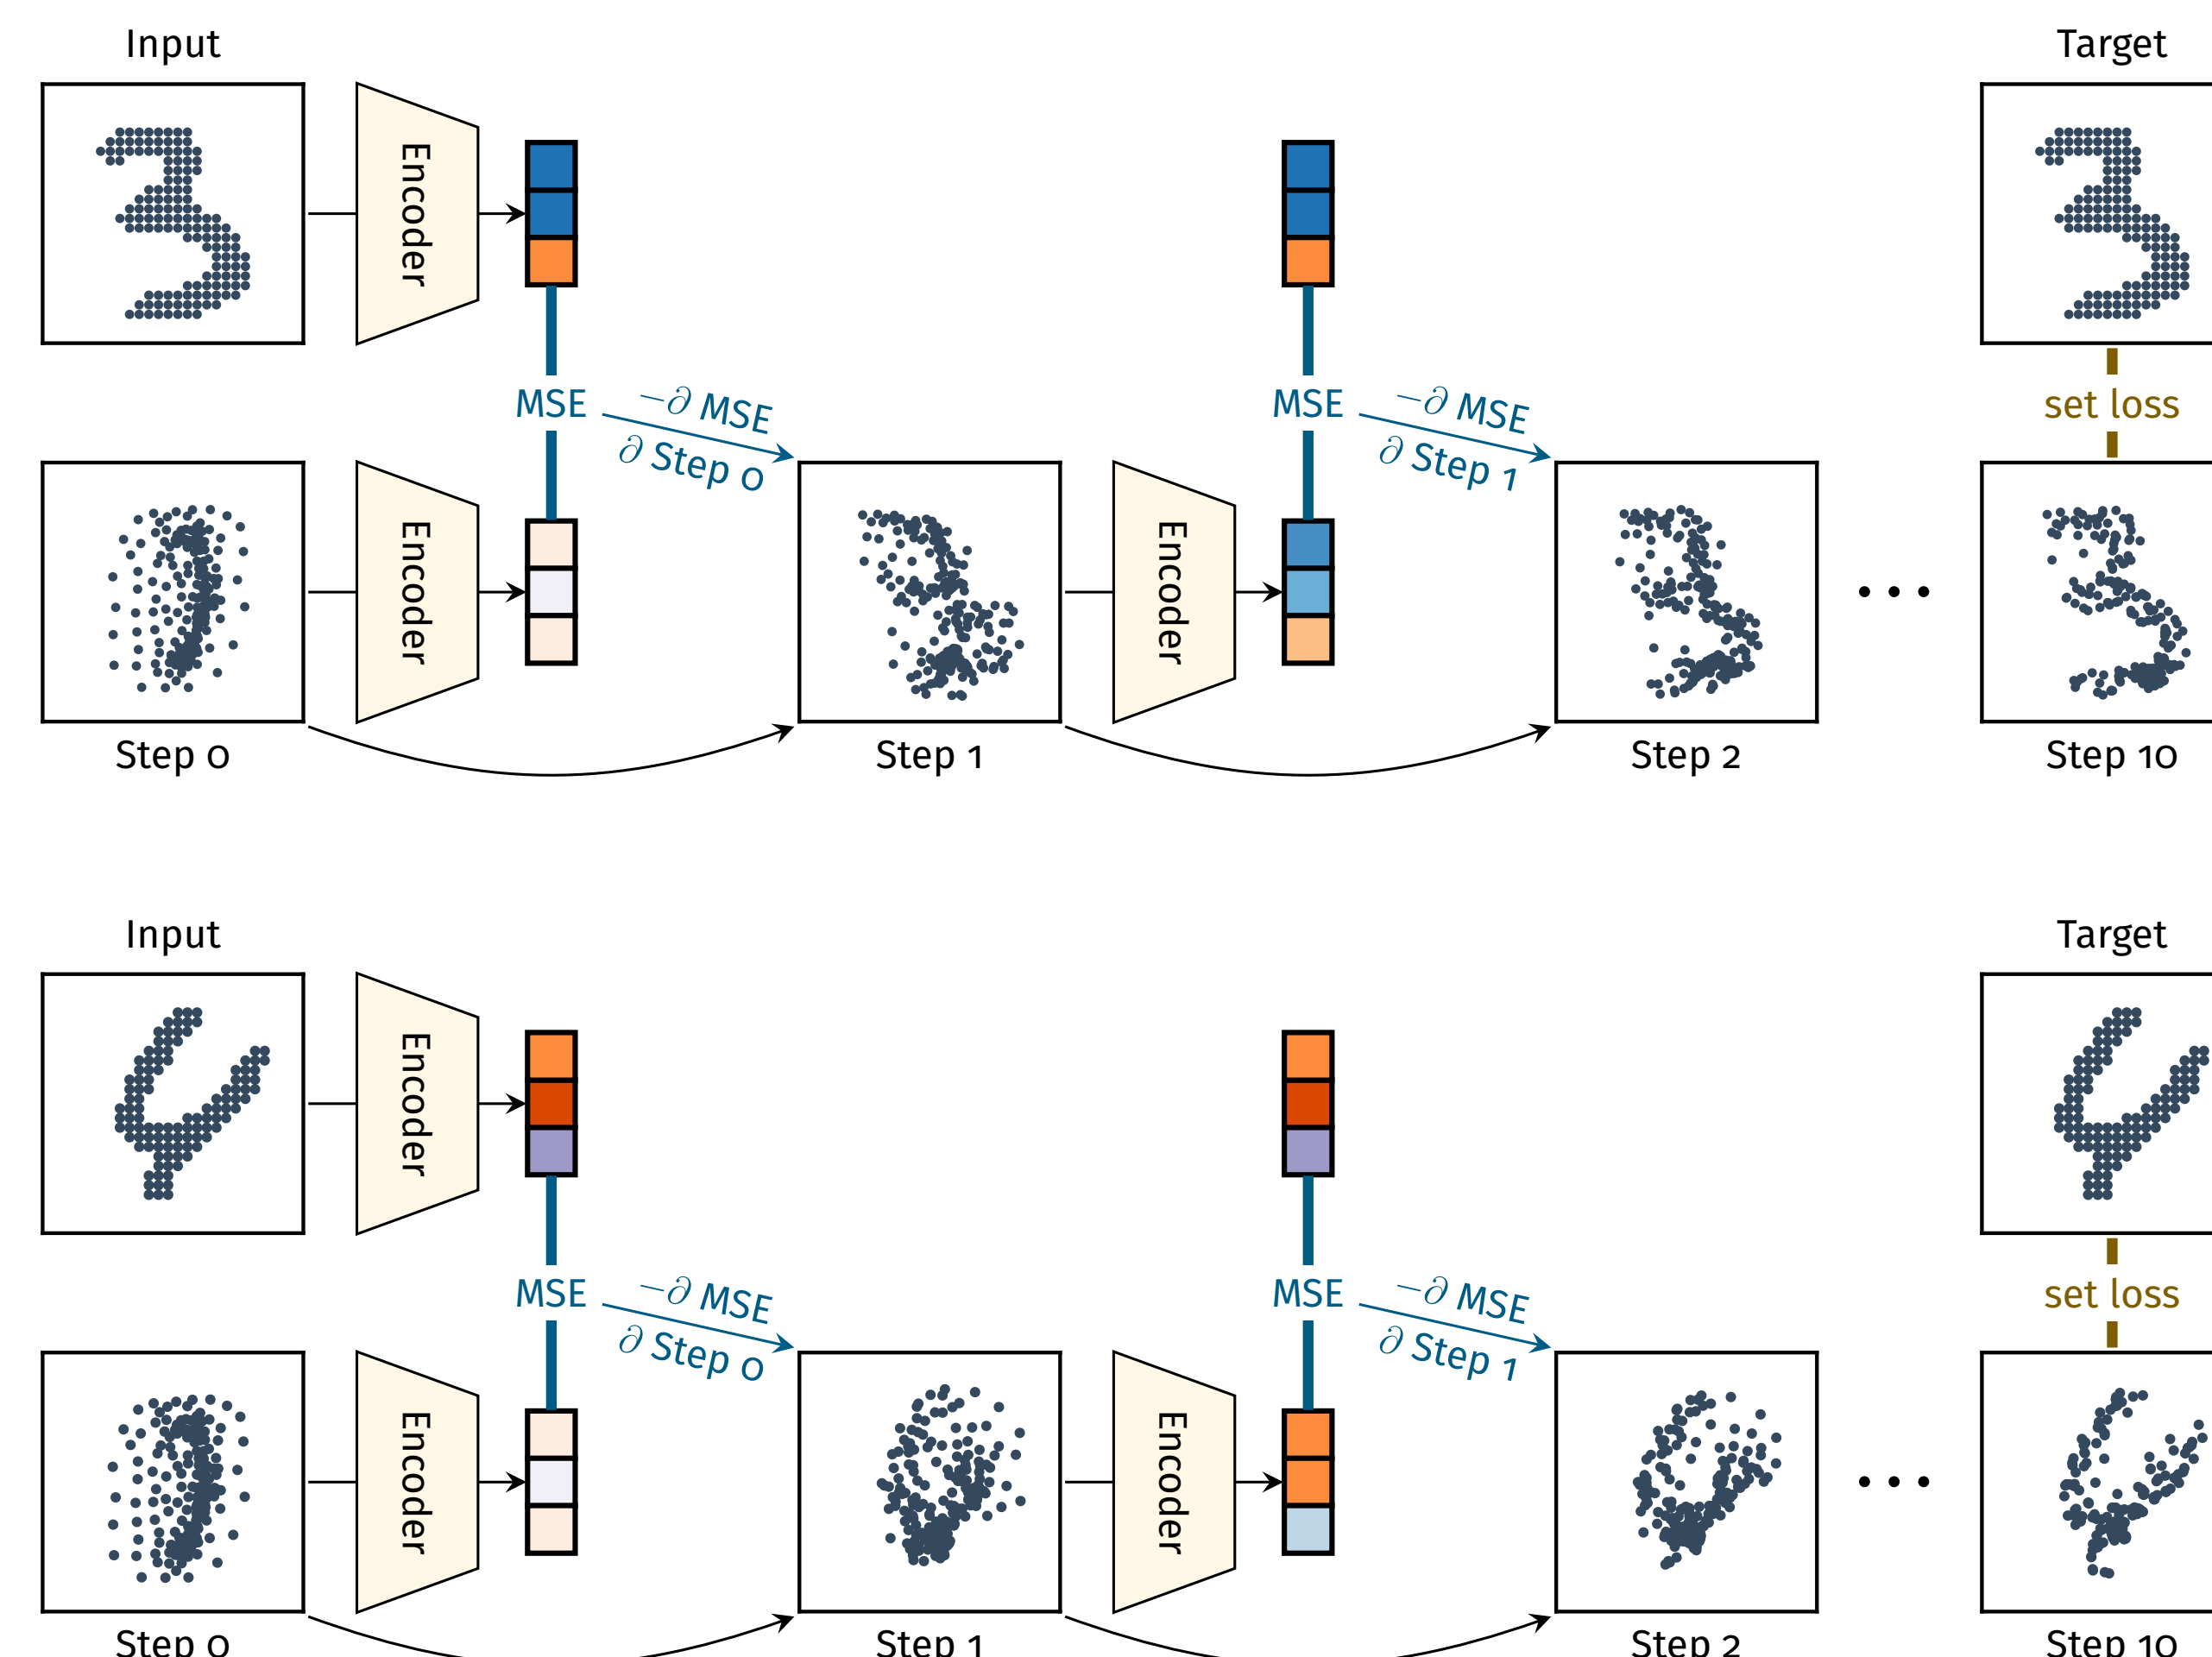
Code and pre-trained models available at
<https://github.com/Cyanogenoid/dspn>



The idea

- Similar set inputs encode to *similar* feature vectors.
- Different set inputs encode to *different* feature vectors.


→ Minimise the difference between predicted and target set by minimising the difference between their feature vectors.




- Train (shared) encoder weights by minimising the **set loss**.
- Gradients of permutation-invariant functions are *equivariant*.
 - All gradient updates $\partial \text{MSE} / \partial \text{set}$ don't rely on the order of the set.
 - Our model is completely **unordered**, exactly what we wanted!

Bounding box set prediction

Bounding box prediction	AP ₅₀	AP ₉₀	AP ₉₅	AP ₉₈	AP ₉₉
MLP baseline	99.3 \pm 0.2	94.0 \pm 1.9	57.9 \pm 7.9	0.7 \pm 0.2	0.0 \pm 0.0
RNN baseline	99.4 \pm 0.2	94.9 \pm 2.0	65.0 \pm 10.3	2.4 \pm 0.0	0.0 \pm 0.0
DSPN (train 10 steps, eval 10 steps)	98.8 \pm 0.3	94.3 \pm 1.5	85.7 \pm 3.0	34.5 \pm 5.7	2.9 \pm 1.2
DSPN (train 10 steps, eval 20 steps)	99.8 \pm 0.0	98.7 \pm 1.1	86.2 \pm 7.2	24.3 \pm 8.0	1.4 \pm 0.9
DSPN (train 10 steps, eval 30 steps)	99.8 \pm 0.1	96.7 \pm 2.4	75.5 \pm 12.3	17.4 \pm 7.7	0.9 \pm 0.7



- Simply replace input encoder with ConvNet image encoder.
- Add **MSE loss** to **set loss** when training the encoder and ResNet weights.
 - Forces minimisation of **MSE** to converge to something sensible.



Object detection

Object attribute prediction	AP _∞	AP ₁	AP _{0.5}	AP _{0.25}	AP _{0.125}
MLP baseline	3.6 \pm 0.5	1.5 \pm 0.4	0.8 \pm 0.3	0.2 \pm 0.1	0.0 \pm 0.0
RNN baseline	4.0 \pm 1.9	1.8 \pm 1.2	0.9 \pm 0.5	0.2 \pm 0.1	0.0 \pm 0.0
DSPN (train 10 steps, eval 10 steps)	72.8 \pm 2.3	59.2 \pm 2.8	39.0 \pm 4.4	12.4 \pm 2.5	1.3 \pm 0.4
DSPN (train 10 steps, eval 20 steps)	84.0 \pm 4.5	80.0 \pm 4.9	57.0 \pm 12.1	16.6 \pm 9.0	1.6 \pm 0.9
DSPN (train 10 steps, eval 30 steps)	85.2 \pm 4.8	81.1 \pm 5.2	47.4 \pm 17.6	10.8 \pm 9.0	0.6 \pm 0.7

Input	Step 5	Step 10	Step 20	Target
	x, y, z = (-0.14, 1.16, 3.57) large purple rubber sphere	x, y, z = (-2.33, -2.41, 0.73) large yellow metal cube	x, y, z = (-2.33, -2.42, 0.78) large yellow metal cube	x, y, z = (-2.42, -2.40, 0.70) large yellow metal cube
	x, y, z = (0.01, 0.12, 3.42) large gray metal cube	x, y, z = (-1.20, 1.27, 0.67) large purple rubber sphere	x, y, z = (-1.21, 1.20, 0.65) large purple rubber sphere	x, y, z = (-1.18, 1.25, 0.70) large purple rubber sphere
	x, y, z = (0.67, 0.65, 3.38) small purple metal cube	x, y, z = (-0.96, 2.54, 0.36) small gray rubber sphere	x, y, z = (-0.96, 2.59, 0.36) small gray rubber sphere	x, y, z = (-1.02, 2.61, 0.35) small gray rubber sphere
	x, y, z = (0.67, 1.14, 2.96) small purple rubber sphere	x, y, z = (1.61, 1.57, 0.36) small yellow metal cube	x, y, z = (1.58, 1.62, 0.38) small purple metal cube	x, y, z = (1.74, 1.53, 0.35) small purple metal cube